

Advice Moments

Growing advice adoption for self-directed Vanguard investors, one well-timed, well-governed moment at a time.

Scope & discretion. Independent thinking exercise built entirely from public information. Not affiliated with, endorsed by, or based on confidential information from Vanguard. Product names are public; figures are illustrative and directional.

WORKING BACKWARDS · THE PRESS RELEASE

ILLUSTRATIVE FUTURE-STATE NARRATIVE

Advice, offered the moment it helps

An illustrative narrative, not a real Vanguard announcement

A self-directed investor changes jobs and her cash balance climbs. The next time she logs in, she sees one quiet, relevant prompt rather than a generic dashboard: a plain-language note that, given her rollover and cash, an advice option might fit, with the fee, the suitability, and a one-tap way to talk to someone or start digitally. She is oriented, not sold to, and books a consultation in under two minutes.

Behind that single prompt is Advice Moments: a governed personalization layer that detects a genuine need, scores whether advice would actually benefit this investor and which tier fits, and delivers a suitability-checked, explainable nudge, routing to a human or a digital flow. Every prompt is capped, reversible, and explainable, and the whole system is measured by whether it moves investors to advice that genuinely suits them, never by clicks alone.

The goal is not to push everyone up the advice staircase. It is to match each investor to the step that fits, including staying self-directed.

FAQ · FOR INVESTORS

Why am I seeing this?

Because a real change in your situation suggested an advice option might help. Every prompt explains its reason, and one tap turns these off for good.

Is this a sales pitch?

No. Prompts are suitability-checked first, show the fee and the fit honestly, and are capped in frequency. If advice would not genuinely help you, you should not see one.

What happens to my data?

Signals are used to decide relevance only, privacy-respecting and opt-out aware, with no prompt sent unless it passes a suitability gate.

FAQ · FOR THE TEAM

What is the wedge?

The activation gap between self-directed and the first advice step, where the most movable assets sit and where the next best step is least often surfaced in-product, at the right moment.

Why will this work where a generic "get advice" banner does not?

Relevance and timing. The nudge is triggered by a real moment, scored for genuine benefit, and written in plain language, not a persistent ad. It is judged on incremental lift against a holdout.

How do we ship without taking on model or compliance risk?

Phase 0 is deterministic rules plus a curated, suitability-gated prompt library, no ML. It proves the surface and earns trust. Phase 1 adds predictive scoring behind evals; Phase 2 adds grounded generative guidance last, behind the strongest guardrails.

How is it governed?

As a behavioral contract: numeric eval thresholds before release, explicit failure modes with detection and recovery, a Reg BI / fiduciary suitability gate, frequency caps, explainability, one-tap opt-out, and EU AI Act-style risk classification and documentation.

What would make this fail?

Over-nudging that erodes trust, or any unsuitable recommendation. Both are designed against directly: caps and suppression for the first, a hard pre-send suitability gate for the second. Guardrail metrics are release gates, not dashboards.

THE ONE METRIC THAT MATTERS

NORTH STAR

Appropriately advised assets (incremental)

Assets moved to a genuinely suitable advice step, measured as true lift against a holdout. Guardrails (suitability, complaints, opt-outs) must not regress, or it does not count as growth.

THREE OPEN QUESTIONS FOR THE TEAM

1. Where does the activation gap leak most today, and what is the real constraint on advice growth: surfacing the moment, advisor capacity, or compliance throughput?
2. How is suitability operationalized today, and what would a pre-send suitability gate need to integrate with?
3. What take-up number does the team already trust, so the evals anchor to reality rather than a guess?

Sources: public information on Vanguard's business, advice continuum, and investor-owned structure; public statements on broadening access to advice; U.S. SEC Regulation Best Interest; the EU AI Act's risk-based approach; current public practice on AI product evals and failure-mode design. Independent work by Nayan Lal. nayanlal1909@gmail.com · linkedin.com/in/nayan-lal